

18Agents in production
across all modules

17Unique LLM endpoints
includes 1 internal-only call

5Product modules
Pipeline · CoPilot · Jobs · Candidates · POFU

3Claude models in use
Opus 4.6 · Sonnet 4.6 · Haiku 4.5

01Pipeline Sessions10 agents

Agentic 7-step hiring drive — Steps 2 through 6 each invoke AI agents.

JD Enhancer

POST /enhance-jd

TRIGGER

Step 2 → recruiter clicks "Generate Assets"

INPUTS

job, clientNotes, companyScript, fields

OUTPUTS

formattedJD · recruiterBrief · clarificationQuestions · reachoutMaterial · sourcingKeywords · marketIntelligence

MODEL

Claude Sonnet 4.6

FIRES

1x per session (shared by all candidates)

Resume Quality Evaluator

POST /sessions/:id/candidates/:cid/evaluate

TRIGGER

Step 3 → auto-fires when candidate is added

INPUTS

resume text, job title, required & preferred skills

OUTPUTS

overall · contentQuality · aiWritingDetection · experienceLevel · verdict · strengths

MODEL

Claude Haiku 4.5

FIRES

1x per candidate added to the session

AI Content Detector

POST /ai/check-ai-content

TRIGGER

Step 3 → recruiter clicks "Run AI Check" (also exposed standalone in Candidate Database)

INPUTS

resume text (100–3000 chars)

OUTPUTS

verdict (human / likely_human / mixed / likely_ai / ai_generated) · confidence · indicators · summary

MODEL

Claude Haiku 4.5

FIRES

On demand · 0–1 per candidate

VI Question Generator

POST /vi/interviews/:id/generate

TRIGGER

Step 5a → recruiter creates a video interview template

INPUTS

job description, desired question count

OUTPUTS

array of { question_text, question_type (technical/behavioral/situational), estimated_time_minutes }

MODEL

Claude Sonnet 4.6

FIRES

1x per interview template (shared across all VI candidates)

VI Evaluator

POST /vi/candidates/:cid/evaluate

TRIGGER

Step 5a → recruiter clicks "Evaluate" on a completed video interview

INPUTS

job description, transcribed candidate responses (Deepgram)

OUTPUTS

overall_score · hiring_recommendation · evaluation_summary · strengths · weaknesses · 5-competency scores · per-question evaluations

MODEL

Claude Sonnet 4.6

FIRES

1x per candidate per interview

MCQ Question Generator

POST /assessments/:id/ai-generate

TRIGGER

Step 5b → recruiter clicks "Generate Questions"

INPUTS

job context, recruiter brief & phone screening Qs (from Step 2), topic & difficulty preference

OUTPUTS

up to 10 MCQs each with { question_text, 4 options, correct_option, explanation, topic, difficulty }

MODEL

Claude Sonnet 4.6

FIRES

1x per MCQ assessment (shared by all invitees)

MCQ Evaluator

(internal) /assessments/take/:token/submit

TRIGGER

Step 5b → fires automatically after candidate submits

INPUTS

assessment title, score breakdown, per-question correctness

OUTPUTS

summary · strengths · gaps · topic_scores · difficulty_breakdown · recommendation · question_analysis

MODEL

Claude Haiku 4.5

FIRES

1x per submission · async (does not block response)

Coding Question Generator

POST /coding-assessments/:id/ai-generate

TRIGGER

Step 5c → recruiter clicks "Generate Questions"

INPUTS

language, count, difficulty, job context

OUTPUTS

array of { title, problem_statement, starter_code, language, question_type (write/fix/complete), difficulty }

MODEL

Claude Sonnet 4.6

FIRES

1x per coding assessment (auto-balances 40/30/30 across question types)

Call QA + Candidate Report

POST /ai/generate-reports

TRIGGER

Step 6 → fires when a recruiter call ends (and from Calling CoPilot)

INPUTS

transcript (speaker/text), role, candidateName, recruiterName, JD excerpt, resume excerpt

OUTPUTS

qaReport (8-dim scorecard, red flags, weak/better nudges) + candidateReport (verdict, scores, recommendation)

MODEL

Claude Opus 4.6

FIRES

1x per completed recruiter call · powers the Recruiter QA module

Coding Evaluator

(internal) /coding-assessments/take/:token/submit

TRIGGER

Step 5c → fires automatically after candidate submits

INPUTS

code solutions per question, language, difficulty, type metadata

OUTPUTS

overall_score · summary · recommendation · strengths · gaps · per-question { correctness, quality, feedback }

MODEL

Claude Haiku 4.5

FIRES

1x per submission · async

02Calling CoPilot3 agents

Real-time AI coaching during live recruiter calls + the auto-demo simulation.

Live Follow-Up Suggester

POST /ai/suggest

TRIGGER

Calling CoPilot → debounced ~1s after each candidate response

INPUTS

last ~6 candidate utterances, current question context

OUTPUTS

2–3 sharp follow-up questions

MODEL

Claude Haiku 4.5

FIRES

Once per new candidate transcript line · low-latency requirement

Candidate Simulation Reply

POST /ai/sim-reply

TRIGGER

Demo Vid / Auto Demo flow — generates a candidate response after each recruiter question

INPUTS

recruiter message, conversation history (last 10 turns), role, candidate name, JD, resume

OUTPUTS

{ candidateReply: 2–4 sentences in candidate persona }

MODEL

Claude Haiku 4.5

FIRES

1x per recruiter question in a simulated call (~8–10 per demo)

Interview Question Generator

POST /ai/generate-questions

TRIGGER

Calling CoPilot → recruiter clicks "Generate Guide" in the Interview Panel

INPUTS

JD (file or text) + resume (optional)

OUTPUTS

4–6 sections (Background · Technical · Problem-Solving · Role-Specific · Culture · Closing) with 3–5 questions each

MODEL

Claude Opus 4.6

FIRES

1x per interview guide build · pre-call preparation

03Job Management2 agents

Hiring-manager intake and per-job JD asset generation.

HM Clarification Question Generator

POST /jobs/:id/qualification-questions

TRIGGER

Job detail → recruiter clicks "Generate Intake Questions" (or regenerates a single Q)

INPUTS

job title, client, required skills, description · (optional) regenerateIndex + existingQuestions

OUTPUTS

7-section intake Q&A structure (or 1 replacement question)

MODEL

Claude Sonnet 4.6

FIRES

1x per job, plus on-demand single-question regeneration

Job-attached JD Refresher

PATCH /jobs/:id/qualify → triggerJDRefresh()

TRIGGER

Recruiter saves HM Q&A and marks the job qualified

INPUTS

job details + HM Q&A answers

OUTPUTS

Refreshed JD assets (formattedJD · recruiterBrief · sourcingKeywords · reachoutMaterial · clarificationQuestions)

MODEL

Claude Sonnet 4.6

FIRES

1x per job qualification · async background refresh

04Candidate Database2 agents

Resume ingestion + AI-content checking on stored profiles.

Resume Parser / Field Extractor

POST /candidates/parse-resume

TRIGGER

Candidate Database → recruiter uploads a PDF or DOCX resume

INPUTS

extracted resume text (via pdf-parse / mammoth)

OUTPUTS

structured fields: name, email, phone, location, current_title, current_company, experience_years, skills, education, work_history

MODEL

Claude Sonnet 4.6

FIRES

1x per resume uploaded

AI Content Detector (standalone)

POST /ai/check-ai-content

TRIGGER

Candidate detail view → "Run AI Check" button (same agent as Pipeline Step 3, exposed here too)

INPUTS

resume text

OUTPUTS

verdict · confidence · indicators · summary

MODEL

Claude Haiku 4.5

FIRES

On demand · per candidate

05Post-Offer Follow-Up (POFU)1 agent

Auto-generated email outreach across the offer-to-Day-1 lifecycle.

POFU Email Composer

server/utils/pofu-engine.js → generateEmail()

TRIGGER

POFU engine selects a trigger (weekly check-in / silence follow-up / DOJ reminder / BGV request / risk escalation / no-response nudge) and asks Claude for a message

INPUTS

candidate name, role, company, DOJ, current state, risk level, trigger reason

OUTPUTS

{ subject, body } — plain-text email signed "The Zeople Team"

MODEL

Claude Haiku 4.5

FIRES

Multiple per candidate over the offer-to-joining window (~3–6 emails typical)